

▼ Task 1 Download the dataset

The Churn_Modelling.csv dataset is downloaded

▼ Task 2 Load the Dataset

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv('/content/Churn_Modelling.csv')
```

df



	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bal
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	83
2	3	15619304	Onio	502	France	Female	42	8	159
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125
...
9995	9996	15606229	Obijaku	771	France	Male	39	5	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57
9997	9998	15584532	Liu	709	France	Female	36	7	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75
9999	10000	15628319	Walker	792	France	Female	28	4	130

10000 rows × 14 columns



▼ Task 3 Perform Visualizations

▼ Univariate Analysis

```
df['Age'].mean()
```

38.9218


```
df['Balance'].median()
```

97198.54000000001

▼ Bivariate Analysis


```
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
sns.barplot(x = df['Exited'] , y = df['Balance']);
```



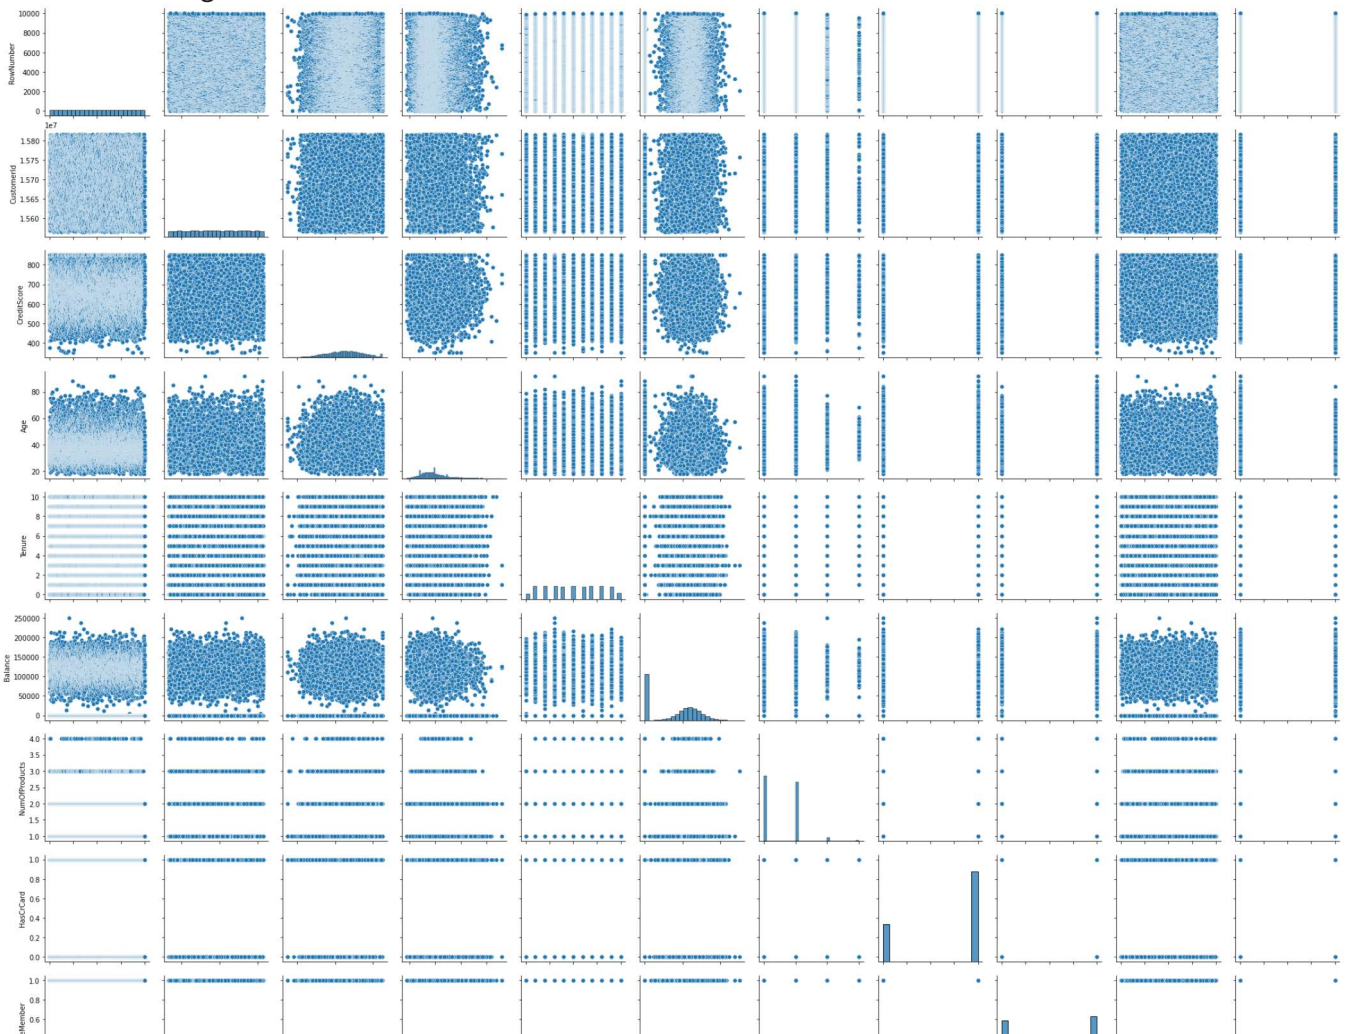
▼ Multi Variate Analysis

```
sns.heatmap(df.corr());
```



```
sns.pairplot(df)
```

<seaborn.axisgrid.PairGrid at 0x7f544f833c50>



Task 4 Perform descriptive statistics on the dataset



```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000

Task 5 Handle the Missing values

```
df.isnull().sum()
```


```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

There are no null values present in the given dataset

Task 6 Find the outliers and replace the outliers



```
sns.boxplot(df['Age']);
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'df': 'df'}. This warning will disappear in Seaborn 0.12.0.





```
sns.boxplot(df['EstimatedSalary']);
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'EstimatedSalary'}. This warning will be removed in a future version of Seaborn.
```




```
sns.boxplot(df['Tenure']);
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'Tenure'}. This warning will be removed in a future version of Seaborn.
```





```
sns.boxplot(df['Balance']);
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'NumOfProducts'}. This warning will be removed in a future version of Seaborn.
```



```
sns.boxplot(df['NumOfProducts']);
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword arguments: {'x': 'NumOfProducts'}. This warning will be removed in a future version of Seaborn.
```



```
outliers=[]  
def detect_outliers(data):  
    threshold=3  
    mean = np.mean(data)  
    std =np.std(data)  
    for i in data:  
        z_score= (i - mean)/std  
        if np.abs(z_score) > threshold:  
            outliers.append(y)  
    return outliers  
outlier_pt=detect_outliers(df)  
outlier_pt
```

Task 7 Check for Categorical columns and perform encoding

```
df.info() #There are few columns that are in object data type instead of int64 or float 64 t
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
df['Geography'] = le.fit_transform(df['Geography'])
```

```
df['Gender'] = le.fit_transform(df['Gender'])
```

```
df
```


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	B
0	1	15634602	Hargrave	619	0	0	42	2	
1	2	15647311	Hill	608	2	0	41	1	83

Task 8 and Task 10 Split the data into dependent and independent variables

```
df.drop(columns = ['RowNumber'])
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num
0	15634602	Hargrave	619	0	0	42	2	0.00	
1	15647311	Hill	608	2	0	41	1	83807.86	
2	15619304	Onio	502	0	0	42	8	159660.80	
3	15701354	Boni	699	0	0	39	1	0.00	
4	15737888	Mitchell	850	2	0	43	2	125510.82	
...	
9995	15606229	Obijiaku	771	0	1	39	5	0.00	
9996	15569892	Johnstone	516	0	1	35	10	57369.61	
9997	15584532	Liu	709	0	0	36	7	0.00	
9998	15682355	Sabbatini	772	1	1	42	3	75075.31	
9999	15628319	Walker	792	0	0	28	4	130142.79	

10000 rows × 13 columns

```
x = df.iloc[:, 0:13].values
y = df.iloc[:, 13:14].values
```

```
from sklearn.model_selection import train_test_split
```

```
xtrain , xtest , ytrain , ytest = train_test_split(x , y , test_size = 0.3 , random_state = 0
```

```
xtrain.shape , xtest.shape
```

```
((7000, 13), (3000, 13))
```

▼ Task 9 Scale the independent variables

```
from sklearn.preprocessing import MinMaxScaler  
from sklearn.preprocessing import StandardScaler
```

```
n = MinMaxScaler()  
s = StandardScaler()
```

```
x = df[['Age', 'Tenure']].values  
y = df['Gender'].values
```

```
n_xtrain = n.fit_transform(xtrain)
```

```
n_xtest = n.fit_transform(xtest)
```